

Inference for Distributions of Categorical Data

lesson 1

Chi-Square Goodness-of-Fit Tests

Chi-Square Goodness-of-Fit Tests

Learning Objectives

After this section, you should be able to...

- ✓ *COMPUTE* expected counts, conditional distributions, and contributions to the chi-square statistic
- ✓ *CHECK* the Random, Large sample size, and Independent conditions before performing a chi-square test
- ✓ *PERFORM* a chi-square goodness-of-fit test to determine whether sample data are consistent with a specified distribution of a categorical variable
- ✓ *EXAMINE* individual components of the chi-square statistic as part of a follow-up analysis

- Here's what the company says about the color distribution of its M&M'S Milk Chocolate Candies: On average, the M&M'S Milk Chocolate Candies will contain 13% of each of browns and reds, 14% yellows, 16% greens, 20% oranges and 24% blues.
- The **one-way table** below summarizes the data from a sample bag of M&M'S Milk Chocolate Candies. In general, one-way tables display the distribution of a categorical variable for the individuals in a sample.

Color	Blue	Orange	Green	Yellow	Red	Brown	Total
Count	9	8	12	15	10	6	60

- Since the company claims that 24% of all M&M'S Milk Chocolate Candies are blue, we could use the one-sample z test for a proportion.

$$H_0: p = 0.24$$

$$H_a: p \neq 0.24$$

where p is the true population proportion of blue M&M'S. We could then perform additional significance tests for each of the remaining colors.

× pretty inefficient ×

■ Comparing Observed and Expected Counts

We need a new kind of significance test, called
chi-square goodness-of-fit test

The null hypothesis in a chi-square goodness-of-fit test should state a claim about the distribution of a single categorical variable in the population of interest. In our example, the appropriate null hypothesis is

H_0 : The company's stated color distribution for
M&M'S Milk Chocolate Candies is correct.

The alternative hypothesis in a chi-square goodness-of-fit test is that the categorical variable does *not* have the specified distribution. In our example, the alternative hypothesis is

H_a : The company's stated color distribution for
M&M'S Milk Chocolate Candies is not correct.

■ Comparing Observed and Expected Counts

We can also write the hypotheses in symbols as:

$$H_0: p_{blue} = 0.24, p_{orange} = 0.20, p_{green} = 0.16, p_{yellow} = 0.14, p_{red} = 0.13, p_{brown} = 0.13,$$

$$H_a: \text{At least one of the } p_i\text{'s is incorrect}$$

where p_{color} is the true **population** proportion of M&M'S Milk Chocolate Candies of that color.

Expected Counts

A sample bag of M&M's milk Chocolate Candies contained 60 candies.

Assume H_0 is correct:

- H_0 : 13% of each of browns and reds, 14% yellows, 16% greens, 20% oranges and 24% blues.
- Calculate the expected counts for each color.

Color	Observed	Expected
Blue	9	14.40
Orange	8	12.00
Green	12	9.60
Yellow	15	8.40
Red	10	7.80
Brown	6	7.80

The idea of the chi-square goodness-of-fit test:

- Assume H_0 is true, compare the **observed counts** (from sample) with the counts that would be expected.
- The more the observed counts differ from the **expected counts**, the more evidence we have against the null hypothesis.

Definition:

The **chi-square statistic** is a measure of how far the observed counts are from the expected counts. The formula for the statistic is

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}}$$

where the sum is over all possible values of the categorical variable.

■ Example: Return of the M&M's

The table shows the observed and expected counts for our sample of 60 M&M's Milk Chocolate Candies. Calculate the chi-square statistic.

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}}$$

Color	Observed	Expected
Blue	9	14.40
Orange	8	12.00
Green	12	9.60
Yellow	15	8.40
Red	10	7.80
Brown	6	7.80

$$\chi^2 = \frac{(9 - 14.40)^2}{14.40} + \frac{(8 - 12.00)^2}{12.00} + \frac{(12 - 9.60)^2}{9.60} + \frac{(15 - 8.40)^2}{8.40} + \frac{(10 - 7.80)^2}{7.80} + \frac{(6 - 7.80)^2}{7.80}$$

$$\chi^2 = 2.025 + 1.333 + 0.600 + 5.186 + 0.621 + 0.415 = 10.180$$

Chi-Square Goodness-of-Fit Tests


■ The Chi-Square Distributions

The Chi-Square Distributions


The chi-square distributions are a family of distributions that take only positive values and are skewed to the right.

The chi-square goodness-of-fit test uses the chi-square distribution with $df = n - 1$. n is the number of categories

Chi-Square Goodness-of-



■ Example: Return of the M&M's



- Since our P-value(0.07) is greater than $\alpha(0.05)$.
- Therefore, we fail to reject H_0 .
- We don't have sufficient evidence to conclude that the company's claimed color distribution is incorrect.

■ Carrying Out a Test

Conditions: Use the chi-square goodness-of-fit test when

✓ **Random**

The data come from a random sample or a randomized experiment.

✓ **Large Sample Size**

All expected counts are at least 5.

✓ **Independent**

Individual observations are independent.

When sampling without replacement, check that the population is at least 10 times as large as the sample (the 10% condition).

The Chi-Square Goodness-of-Fit Test

Suppose the k conditions are met. To determine if the distribution, expressed in terms of category, is a good fit to the observed proportions, we use the chi-square goodness-of-fit test.

Before we start using the chi-square goodness-of-fit test, we have two important cautions to offer.

1. The chi-square test statistic compares observed and expected *counts*. Don't try to perform calculations with the observed and expected *proportions* in each category.
2. When checking the Large Sample Size condition, be sure to examine the *expected* counts, not the observed counts.

where the sum is over all categories. The p -value is the area to the right of χ^2 under the density curve of the chi-square distribution with $k - 1$ degrees of freedom.



■ Example: When Were You Born?

Are births evenly distributed across the days of the week? The one-way table below shows the distribution of births across the days of the week in a random sample of 140 births from local records in a large city. Do these data give significant evidence that local births are not equally likely on all days of the week?

Day	Sun	Mon	Tue	Wed	Thu	Fri	Sat
Births	13	23	24	20	27	18	15

State: We want to perform a test of

H_0 : Birth days in this local area are evenly distributed across the days of the week.

H_a : Birth days in this local area are not evenly distributed across the days of the week.

The null hypothesis says that the proportions of births are the same on all days. In that case, all 7 proportions must be $1/7$. So we could also write the hypotheses as

$$H_0: p_{Sun} = p_{Mon} = p_{Tues} = \dots = p_{Sat} = 1/7.$$

H_a : At least one of the proportions is not $1/7$.

We will use $\alpha = 0.05$.

Plan: If the conditions are met, we should conduct a chi-square goodness-of-fit test.

- *Random* The data came from a random sample of local births.
- *Large Sample Size* Assuming H_0 is true, we would expect one-seventh of the births to occur on each day of the week. For the sample of 140 births, the expected count for all 7 days would be $1/7(140) = 20$ births. Since $20 \geq 5$, this condition is met.
- *Independent* Individual births in the random sample should occur independently (assuming no twins). Because we are sampling without replacement, there need to be at least $10(140) = 1400$ births in the local area. This should be the case in a large city.

■ Example: When Were You Born?

Do: Since the conditions are satisfied, we can perform a chi-square goodness-of-fit test. We begin by calculating the test statistic.


Test statistic:

$$\begin{aligned}\chi^2 &= \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}} \\ &= \frac{(13-20)^2}{20} + \frac{(23-20)^2}{20} + \frac{(24-20)^2}{20} + \frac{(20-20)^2}{20} \\ &\quad + \frac{(27-20)^2}{20} + \frac{(18-20)^2}{20} + \frac{(15-20)^2}{20} \\ &= 2.45 + 0.45 + 0.80 + 0.00 + 2.45 + 0.20 + 1.25 \\ &= 7.60\end{aligned}$$

P-Value:

Using Table C: $\chi^2 = 7.60$ is less than the smallest entry in the $df = 6$ row, which corresponds to tail area 0.25. The P -value is therefore greater than 0.25.

Using technology: We can find the exact P -value with a calculator: $\chi^2\text{cdf}(7.60, 1000, 6) = 0.269$.



Conclude: Because the P -value, 0.269, is greater than $\alpha = 0.05$, we fail to reject H_0 . These 140 births don't provide enough evidence to say that all local births in this area are not evenly distributed across the days of the week.

Inference for Distributions of Categorical Data

lesson 2

Chi-Square Test for Homogeneity

■ Example

On the campus of a large university, the students are freshmen, sophomores, juniors and seniors. The school officials had heard several complaints about the food services, but felt that the complaints differed across the four grades. To get students input, a random sample of students was selected from each grade. Each selected students was asked, “Are the current food services satisfactory?” The answers were recorded as “satisfactory” or “not satisfactory”. The answers are summarized as follows.

Is there any significant evidence to indicate whether the proportion of students satisfied with the current food services differs in different grades at the significance level of 0.05?

	Satisfactory	Not Satisfactory	Sample size
Freshman	156	44	200
Sophomore	144	56	200
Junior	98	52	150
Senior	88	62	150
Total	486	214	700

■ Example

	Satisfactory	Not Satisfactory	total
Answers	486	214	700

H0: $p_1=0.65$

$p_2=0.35$

Ha: At least one of these proportions are incorrect

p_1 and p_2 are the proportions of Satisfactory and Not Satisfactory respectively

Chi-square goodness-of-fit test

Based on this table, we have that the distribution of current food services satisfactory is:

$P(\text{satis}) = 0.6943$

$P(\text{not satis}) = 0.3057$

Chi-square test for homogeneity

	Satisfactory	Not Satisfactory	Sample size
Freshman	156	44	200
Sophomore	144	56	200
Junior	98	52	150
Senior	88	62	150
Total	486	214	700

$$P(\text{satis}) = 0.6943$$

$$P(\text{not satis}) = 0.3057$$



H₀: The distributions of students satisfied with the current food services are the same across all four grades.

H_a: The distributions of students satisfied with the current food services are the same across all four grades.

■ Example: Does Music Influence Purchases?

Researchers suspect that background music may lead to different buying behaviors. One study in a supermarket compared three randomly assigned treatments: no music, French accordion music, and Italian string music.

Under each condition, the researchers recorded the numbers of bottles of French, Italian, and other wine purchased. The table below summarizes the data.

Wine	Music			Total
	None	French	Italian	
French	30	39	30	99
Italian	11	1	19	31
Other	43	35	35	113
Total	84	75	84	243

H_0 : There is no difference in the distributions of wine purchases at this store when no music, French accordion music, or Italian string music is played.

H_a : There is a difference in the distributions of wine purchases at this store when no music, French accordion music, or Italian string music is played.

■ Finding Expected Counts

Finding Expected Counts

The expected count in any cell of a two-way table **when H_0 is true** is

$$\text{expected count} = \frac{\text{row total} \cdot \text{column total}}{\text{table total}}$$

Expected Counts				
Wine	Music			Total
	None	French	Italian	
French	34.22	30.56	34.22	99
Italian	10.72	9.57	10.72	31
Other	39.06	34.88	39.06	113
Total	84	75	84	243

■ Calculating the Chi-Square Statistic

In order to calculate a chi-square statistic for the wine example, we must check to make sure the conditions are met:

- ✓ All the **expected** counts in the music and wine study are at least **5**.
- ✓ The treatments were assigned at random.

Just as we did with the chi-square goodness-of-fit test, we compare the observed counts with the expected counts using the statistic

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}}$$

■ Calculating The Chi-Square Statistic

Observed Counts					Expected Counts				
Wine	Music			Total	Wine	Music			Total
	None	French	Italian			None	French	Italian	
French	30	39	30	99	French	34.22	30.56	34.22	99
Italian	11	1	19	31	Italian	10.72	9.57	10.72	31
Other	43	35	35	113	Other	39.06	34.88	39.06	113
Total	84	75	84	243	Total	84	75	84	243

The χ^2 statistic is the sum of nine such terms:

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}} = \frac{(30 - 34.22)^2}{34.22} + \frac{(39 - 30.56)^2}{30.56} + \dots + \frac{(35 - 39.06)^2}{39.06}$$

$$= 0.52 + 2.33 + \dots + 0.42 = 18.28$$

■ Example: Does Music Influence Purchases?

P-value = $P(\chi^2 > 18.28) = \dots$, where $\chi^2 \sim \chi^2_{df} = \chi^2_4$.

The small P -value gives us convincing evidence to reject H_0 and conclude that there is a difference in the distributions of wine purchases at this store when no music, French accordion music, or Italian string music is played. Furthermore, the random assignment allows us to say that the difference is caused by the music that's played.

The appropriate hypotheses for a **chi**-square test for homogeneity are:

H_0 : There is no difference in distributions of a categorical variable across populations or treatments.

H_a : There is a difference in distributions of a categorical variable across populations or treatments.

VAR-8.1.2

The appropriate hypotheses for a **chi**-square test for independence are:

H_0 : There is no association between two categorical variables in a given population or the two categorical variables are independent.

H_a : Two categorical variables in a population are associated or dependent.

■ Example: Cell-Only Telephone Users

Random digit dialing telephone surveys used to exclude cell phone numbers. If the opinions of people who have only cell phones differ from those of people who have landline service, the poll results may not represent the entire adult population. The Pew Research Center interviewed separate random samples of cell-only and landline telephone users who were less than 30 years old. Here's what the Pew survey found about how these people describe their political party affiliation.

	Cell-only sample	Landline sample
Democrat or lean Democratic	49	47
Refuse to lean either way	15	27
Republican or lean Republican	32	30
Total	96	104

State: We want to perform a test of

H_0 : There is no difference in the distribution of party affiliation in the cell-only and landline populations.

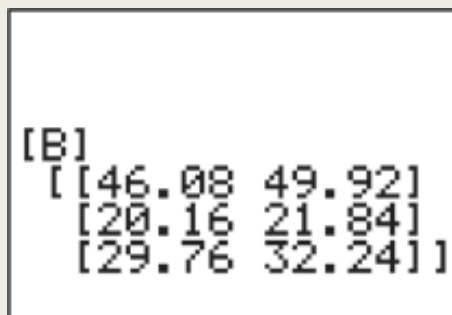
H_a : There is a difference in the distribution of party affiliation in the cell-only and landline populations.

We will use $\alpha = 0.05$.

■ Example: Cell-Only Telephone Users

Plan: If the conditions are met, we should conduct a chi-square test for homogeneity.

- *Random* The data came from separate random samples of 96 cell-only and 104 landline users.
- *Large Sample Size* We followed the steps in the Technology Corner (page 705) to get the expected counts. The calculator screenshot confirms all expected counts ≥ 5 .



```
[B]
[[46.08 49.92]
 [20.16 21.84]
 [29.76 32.24]]
```

- *Independent* Researchers took independent samples of cell-only and landline phone users. Sampling without replacement was used, so there need to be at least $10(96) = 960$ cell-only users under age 30 and at least $10(104) = 1040$ landline users under age 30. This is safe to assume.

■ Example: Cell-Only Telephone Users


Do: Since the conditions are satisfied, we can perform a chi-test for homogeneity. We begin by calculating the test statistic.

Test statistic:

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}}$$
$$= \frac{(49 - 46.08)^2}{46.08} + \frac{(47 - 49.92)^2}{49.92} + \dots + \frac{(30 - 32.24)^2}{32.24} = 3.22$$

P-Value:

Using $df = (3 - 1)(2 - 1) = 2$, the P -value is 0.20.



Conclude: Because the P -value, 0.20, is greater than $\alpha = 0.05$, we fail to reject H_0 . There is not enough evidence to conclude that the distribution of party affiliation differs in the cell-only and landline user populations.

The article “Daily Weigh-ins Can Help You Keep Off Lost Pounds, Experts Say” describes an experiment in which 291 people who had lost at least 10% of their body weight in a medical weight loss program were assigned at random to one of three groups for follow-up. One group met monthly in person, one group “met” online monthly in a chat room, and one group received a monthly newsletter by mail. After 18 months, participants in each group were classified according to whether or not they had regained more than 5 pounds, resulting in the data below.

	AMOUNT OF WEIGHT GAINED		Row Marginal Total
	Regained 5 Lb or Less	Regained More Than 5 Lb	
In-Person	52	45	97
Online	44	53	97
Newsletter	27	70	97

Does there appear to be a difference in the weight regained proportions for the three follow-up methods?

H_0 : The proportions for the two weight-regained categories are the same for the three follow-up methods.

H_a : The weight-regained category proportions are not the same for all three follow-up methods.

Significance level: $\alpha = .01$

Test statistic:
$$X^2 = \sum_{\text{all cells}} \frac{(\text{observed cell count} - \text{expected cell count})^2}{\text{expected cell count}}$$

We will carry out a Chi-square test for the homogeneity.

$\chi^2 \sim \chi^2_{df} = \chi^2_2$ if the following conditions are satisfied:

1. All of the computed expected counts are greater than 5.
2. The subjects in this experiment were assigned at random to the treatment groups.

Calculations:

test statistic $\chi^2 = 13.773$

p-value = $P(\chi^2_2 > 13.773) = 0.001$

Conclusion:

Since P-value(0.001) is less than the significance level of 0.05, we reject H_0 . The weight regained proportions are not the same for the three follow-up methods.

Inference for Distributions of Categorical Data

lesson 3

Chi-Square Test for Independence

■ Relationships Between Two Categorical Variables

A study followed a random sample of 8474 people with normal blood pressure for about four years. All the individuals were free of heart disease at the beginning of the study. Each person took the Spielberg Trait Anger Scale test, which measures how prone a person is to sudden anger. Researchers also recorded whether each individual developed coronary heart disease (CHD). This includes people who had heart attacks and those who needed medical treatment for heart disease. Here is a two-way table that summarizes the data:


	Low anger	Moderate anger	High anger	Total
CHD	53	110	27	190
No CHD	3057	4621	606	8284
Total	3110	4731	633	8474

	Low anger	Moderate anger	High anger	Total
CHD	53	110	27	190
No CHD	3057	4621	606	8284
Total	3110	4731	633	8474

Find the conditional probability of CHD given different levels of anger.

There is a clear trend:
as the anger score increases, so does the percent who suffer heart disease.

Do the data provide convincing evidence of an **association** between anger level and heart disease in the population of interest?



■ The Chi-Square Test for Association/Independence

H_0 : There is **no** association between anger level and heart disease in the population of people with normal blood pressure.

H_a : There is an association between anger level and heart disease in the population of people with normal blood pressure.

No association between two variables means the variables are independent.

H_0 : Anger and heart disease are independent in the population of people with normal blood pressure.

H_a : Anger and heart disease are **not** independent in the population of people with normal blood pressure.

■ Example: Angry People and Heart Disease

	Observed			Expected		
	Low	Moderate	High	Low	Moderate	High
CHD	53	110	27	69.73	106.08	14.19
No CHD	3057	4621	606	3040.27	4624.92	618.81

State: We want to perform a test of

H_0 : There is no association between anger level and heart disease in the population of people with normal blood pressure.

H_a : There is an association between anger level and heart disease in the population of people with normal blood pressure.

We will use $\alpha = 0.05$.

■ Example: Angry People and Heart Disease

Plan:

If the conditions are met, we should conduct a chi-square test for independence.

- *Random*

The data came from a random sample of 8474 people with normal blood pressure.

- *Large Sample Size*

All the expected counts are at least 5, so this condition is met.

- *Independent*

Because we are sampling without replacement, we need to check that the total number of people in the population with normal blood pressure is at least $10(8474) = 84,740$. This seems reasonable to assume.

Do: Since the conditions are satisfied, we can perform a chi-test for association/independence. We begin by calculating the test statistic.

Test statistic :

$$\begin{aligned}\chi^2 &= \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}} \\ &= \frac{(53 - 69.73)^2}{69.73} + \frac{(110 - 106.08)^2}{106.08} + \dots + \frac{(606 - 618.81)^2}{618.81} \\ &= 4.014 + 0.145 + \dots + 0.265 = 16.077\end{aligned}$$

P-Value:

The two-way table of anger level versus heart disease has 2 rows and 3 columns. We will use the chi-square distribution with $df = (2 - 1)(3 - 1) = 2$ to find the P -value.

$P\text{-value} = P(\chi^2 > 16.077) = 0.00032$, where $\chi^2 \sim \chi^2_2$

Conclude: Because the P -value (0.00032) is less than $\alpha = 0.05$, we reject H_0 and conclude that anger level and heart disease are associated in the population of people with normal blood pressure.

<http://www.ltcconline.net/greenL/java/statistics/catstatprob/categorizingstatproblemsjavascript.html>